

# Introduction to neural networks

Laboratory of robotics and control systems  
ITT PROJECT MANAGEMENT SERVICES L.L.C

# What is this course about?

**Aim:** provide students with basic skills in the field of deep learning computer vision in the context of robotics

**Focus:**

1. process of integrating and running neural nets on robots
2. optimizing model's performance and improving their quality
3. choosing model architecture for solving specific tasks related to the robot's perception of the environment

**Caution:** This course is more practical than fundamental. That's why some historical details, proofs, overview of some tasks and topics, not related to robotics, might be omitted

# Course Structure

1. Intro to DL, Neuron model, Perceptron, Losses, Backpropagation, PyTorch concpetions
2. Activation functions, Optimizers, Dataset splitting, Underfitting, Overfitting, Metrics
3. CV problems and data, Convolutional neural nets, More Pytorch Concepts, GPU
4. NNs architecutres: AlexNet, VGG, ResNet, Inception, MovableNet. Pretraining and Finetuning theory
5. Semantic segmentation: data, encoder-decoder, U-Net architecture, losses, metrics. SAM. Instance Segmentation
6. Object detection: data, metrics. Yolo3-v8 architecture
7. Object tracking
8. Pose Estimation
9. Point Clouds
10. Software integration: LibTorch, TorchScript, ONNX, OpenVINO. SAHR Vision Pipeline
11. Quantization
12. Pruning
13. Domain Adaptation, Self-supervised learning (DINO)

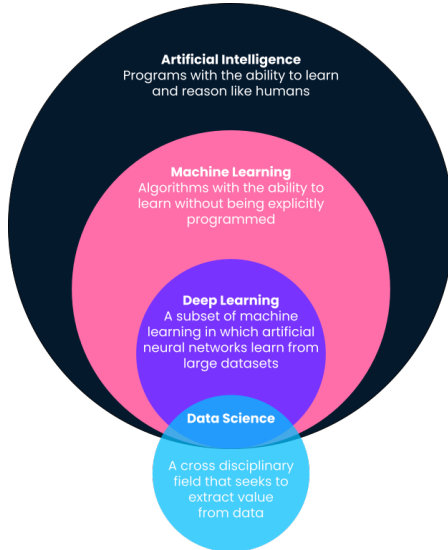
# Course prerequisites

1. Basic Python, Numpy
2. Basic linear algebra, probability theory
3. CVR (Classic part) recommended

# Plan

1. What is Machine learning?
2. Deep learning
3. Perceptron
4. Training process
5. PyTorch

# Relation between AI, ML, DL



# Types of Machine learning

Machine learning divides into:

- Supervised learning (have answers)
- Unsupervised learning (have no answers)
- Reinforcement learning (agent, environment, reward)
- Semi-supervised learning (partial labeling)
- Self-supervised learning (no answers, synthetic task in order to obtain good object representation)

# Supervised learning

Provided:

- $X$  - set of objects
- $Y$  - set of answers
- $y : X \rightarrow Y$  - unknown dependency

We have to find  $a : X \rightarrow Y$  that maps object  $x$  to expected output  $y$ .

Problems:

- Classification
- Regression
- Object detection
- Semantic segmentation



# Regression vs Classification part 1

## Regression



What will be the temperature tomorrow?

84°



Fahrenheit

## Classification



Will it be hot or cold tomorrow?


COLD

HOT




Fahrenheit

## Regression vs Classification part 2



**Classification**

Given 2 features  $(x, y)$  predict  
dividing surface



**Regression**

Given 1 feature  $x$  predict  $y$

# Unsupervised learning


## Features:

- Learns patterns exclusively from unlabeled data
- Due to missing labels, quality assessment is non-trivial

## Problems:


- Clustering
- Dimension reduction
- Anomaly detection

# Clustering




Given unlabeled data with 2 features ( $x$ ,  $y$ ), split objects into groups

# Dimensionality reduction




Remove excess information from data by projecting it into a lower dimensional space

# Anomaly detection



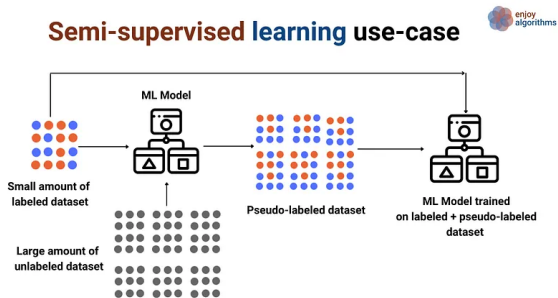
Find samples, that came from another data distribution

# Reinforcement learning



- Agent performs action in environment
- Environment responds by giving it's new state and reward for action
- Environment serves as a teacher for the agent

# Semi-supervised learning

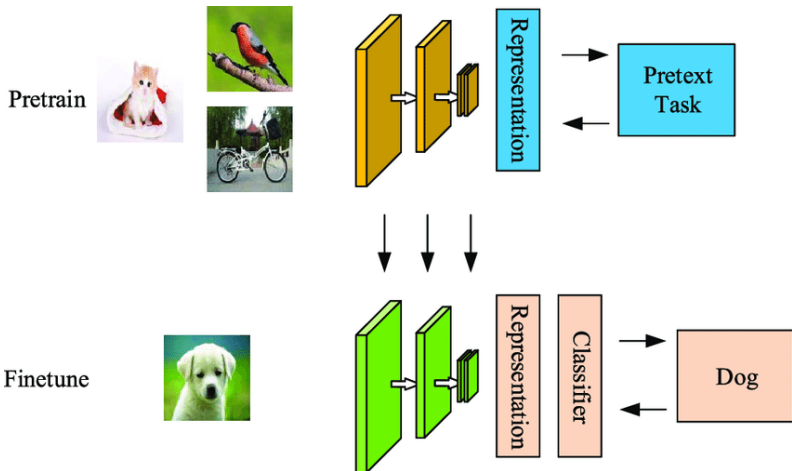


Iterative process:

1. Label small amount of data
2. Train base model on labeled data
3. Generate pseudo-labels for unlabeled data
4. Take the most confident predictions
5. Add them to labeled partition
6. Repeat 2. - 6.



# Self-supervised learning



# Table of Contents

1. What is Machine learning?

2. Deep learning


3. Perceptron

4. Training process


5. PyTorch

# Classic ML vs Deep learning


## Machine Learning





## Deep Learning



# Comparison of scaling laws




# Neuron structure



# Linear layer

Input layer

Output layer



A simple neural network

$$\begin{bmatrix} w_1 & w_2 & w_3 & w_4 \\ w_1 & w_2 & w_3 & w_4 \\ w_1 & w_2 & w_3 & w_4 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \end{bmatrix} + \begin{bmatrix} b \\ b \\ b \end{bmatrix} = \begin{bmatrix} w_1x_1 + w_2x_2 + w_3x_3 + w_4x_4 + b \\ w_1x_1 + w_2x_2 + w_3x_3 + w_4x_4 + b \\ w_1x_1 + w_2x_2 + w_3x_3 + w_4x_4 + b \end{bmatrix} \xrightarrow{\text{activation}} \begin{bmatrix} a_1 \\ a_2 \\ a_3 \end{bmatrix}$$

# Table of Contents

1. What is Machine learning?

2. Deep learning


3. Perceptron

4. Training process

5. PyTorch

# Multilayer perceptron (MLP)

Illustrative example of Multilayer perceptron, a Feedforward neural network




$x_1, x_2$  : input data features  
 $w_{ij}$  : weights of the network  
 $h_1, h_2$  : nodes in the hidden layer  
 $y_k$  : output variable

© AIML.com Research

This example consists of: (a) an input layer with two input nodes and a bias node, (b) one hidden layer with two neurons, and (c) an output layer with one neuron



# Some activation functions



# Approximation theorems

Multilayer perceptron with 1 hidden layer and:

- arbitrary width, and sigmoid function (Cybenko, 1989)
- arbitrary width, and any activation function (Hornik, 1991)

may approximate any continuous function with predefined precision  $\varepsilon$

## Weaknesses:

1. Theorems don't answer what is the width of MLP
2. Theorems don't answer how to find the parameters

# Table of Contents

1. What is Machine learning?


2. Deep learning

3. Perceptron

4. Training process

5. PyTorch

# Forward Propagation



# Loss function

- Estimates the quality of model's predictions
- Takes the model's prediction and the correct answer as an input
- Depends on the data type and the problem
- Is being minimized in the training process

# Loss function examples

## Regression:

Mean Squared Error:  $\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$

Mean Average Error:  $\frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i|$


Mean absolute percentage error:  $\frac{1}{n} \sum_{i=1}^n \left| \frac{y_i - \hat{y}_i}{y_i} \right|$

## Binary Classification:

Cross Entropy Loss:  $\frac{1}{n} \sum_{i=1}^n -(y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i))$   $(\hat{y}_i \in \{0, 1\})$

Hinge Loss:  $\frac{1}{n} \sum_{i=1}^n \max(1 - y_i \hat{y}_i)$   $(\hat{y}_i \in \{-1, 1\})$

# Gradient descent



$$F(\omega) = \sum_i^N \frac{L(x_i, y_i, \omega)}{N} + \lambda R(\omega)$$

$$\nabla f(\omega_1, \dots, \omega_n) = \left( \frac{\partial f}{\partial \omega_1}, \dots, \frac{\partial f}{\partial \omega_n} \right)$$


$$w := w - \alpha \nabla f$$

# Learning rate

$$w := w - \alpha \nabla f$$


$\alpha$  - learning rate, controls convergence speed

## Gradient Descent

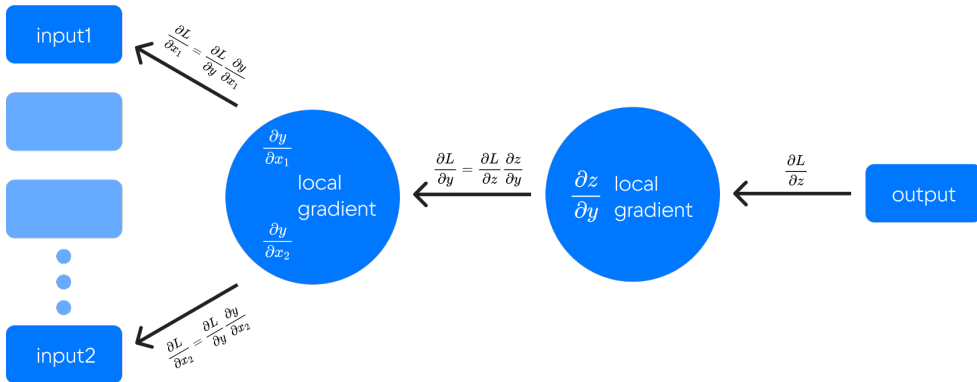




# Loss function of 2 variables example



# Backpropagation



# Backpropagation

$$L = L(O_n(\dots(O_i(O_{i-1}(\dots(O_1))))))$$

$$O_i = O_i(W_i, O_{i-1})$$

$$\frac{\partial L}{\partial W_i} = \frac{\partial L}{\partial O_i} \frac{\partial O_i}{\partial W_i}$$

$$\frac{\partial L}{\partial O_{i-1}} = \frac{\partial L}{\partial O_i} \frac{\partial O_i}{\partial O_{i-1}}$$

$$O_n = Y_{pred}$$

$$O_1 = X$$

# Table of Contents

1. What is Machine learning?

2. Deep learning

3. Perceptron

4. Training process

5. PyTorch

# Pytorch

## Features:

- Automatic differentiation for building and training neural networks
- Native GPU and CPU usage
- Numpy-like Tensors
- Dynamic Computational graph
- Prototyping of neural networks using torch.nn and torchvision

**PyTorch = Numpy + GPU + Autograd**

# Autograd

PyTorch's automatic differentiation engine that powers neural network training

`.backward()`

- computes the gradients from each `.grad_fn`
- accumulates them in the respective tensor's `.grad` attribute
- using the chain rule, propagates all the way to the leaf tensors

# Autograd




Figure 1: Example of an augmented computational graph



# Thank you for your time

Laboratory of robotics and control systems  
ITT PROJECT MANAGEMENT SERVICES L.L.C